Neo4J: Sample Questions 3

Contents

Question 32 of 80

Question 33 of 80

Question 34 of 80

Question 35 of 80

Question 36 of 80

Question 37 of 80

Question 38 of 80

Question 39 of 80

Question 40 of 80

Question 41 of 80

Question 42 of 80

Question 43 of 80

Question 44 of 80

Question 45 of 80

Question 46 of 80

Question 47 of 80

Question 48 of 80

Question 49 of 80

Question 50 of 80

Question 51 of 80

Question 52 of 80

Question 53 of 80

Question 54 of 80

Question 55 of 80

Question 56 of 80

Question 57 of 80

Question 58 of 80

Question 59 of 80

Question 60 of 80

Question 61 of 80

Question 62 of 80

Question 63 of 80

Question 64 of 80

Question 65 of 80

Question 66 of 80

Question 67 of 80

Question 68 of 80

Question 69 of 80

Question 70 of 80

Question 71 of 80

Question 1 of 80

What is the main responsibility of read replicas in a causal cluster?

- A) Scaling writes
- B) Taking part in the Raft election algorithm
- **C) Scaling out graph workloads**
- D) Replicating data around the cluster

```
Points: 0 out of 1
Correct answer: C)
```

Question 2 of 80

Which of the following would be the correct Cypher syntax to create an index on the id property for nodes with the Customer label?

- **A) CREATE INDEX ON :Customer(id);**
- B) CREATE INDEX ON :Customer.id;
- C) CREATE INDEX ON Customer(id);
- D) CREATE GRAPH INDEX ON :Customer.id;

```
Points: 0 out of 1
Correct answer: A)
```

Question 3 of 80

You can create the same type of relationship between two nodes with each relationship having different properties?

- A) **True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 4 of 80

Which of the following Cypher statements will return actors and the directors who directed their movies?

- **A)**

```
MATCH (actor)–[:ACTED_IN]–>(movie)<–[:DIRECTED]–(director)
RETURN actor, director
```

- B)

```
MATCH (actor)–[:ACTED_IN]–>(movie)
JOIN (movie)<–[:DIRECTED]–(director)
RETURN actor, director
```

- C)

```
MATCH (actor)–[:ACTED_IN]–(movie)
CONNECT (movie)–[d:DIRECTED]–(director)
RETURN actor, director
```

- D)

```
MATCH (actor)–[a:ACTED_IN]–>(movie)<–[b:DIRECTED]–(director)
RETURN a, b
```

```
Points: 0 out of 1
Correct answer: A)
```

Question 5 of 80

Which of the following are options for querying data from Neo4j in version 3.2 and later?

- **A) Through the Bolt protocol available with the server.**
- **B) Using the Cypher shell utility.**
- C) Through the message queue service that ships with the server.
- D) Through the JDBC driver supplied by Neo4j Inc. and shipped with the server.
- **E) Using the Neo4j Browser web interface that ships with Neo4j.**

```
Points: 1 out of 1
Correct answer: A) B) E)
```

Question 6 of 80

What is the result of the following statement: RETURN null <> false

- A) true
- B) false
- **C) null**

- D) none of the above

```
Points: 0 out of 1
Correct answer: C)
```

Question 7 of 80

A single node can have a relationship that points at itself?

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 8 of 80

What, specifically, does the following expression do?: size((n)-[:TYPE]->()) ?

- **A) It counts the number of relationships with this type and direction?**
- B) This is not valid Cypher.
- **C) It reads the degree-value for this relationship-type and direction directly from the node record.**
- D) It executes a MATCH first and then returns the size of the path - collection.

```
Points: 0 out of 2
Correct answer: A) C)
```

Question 9 of 80

How does the Neo4j configuration property "keep_logical_logs" affect Neo4j backup?

- **A) This property defines how long to keep the log files with transaction history. When running a backup, if the time since you've last run a backup is longer than the time specified in this configuration setting, Neo4j will be unable to incrementally apply all historical transactions and instead will do a full backup.**
- B) This configuration property identifies the number of logs to keep to recover a backup into a fully functioning instance.
- C) This configuration will define whether logical or physical logs will be stored by Neo4j, which will determine whether a backup can copy data from memory or from the disk.
- D) This configuration identifies how many logs Neo4j should keep in memory for quick, incremental backups every hour.

```
Points: 0 out of 1
Correct answer: A)
```

Question 10 of 80

How many nodes can a single relationship connect?

- A) Exactly two.

- B) Exactly one.
- **C) One or two.**
- D) More than two.
- E) None of the above.

```
Points: 0 out of 1
Correct answer: C)
A node can have a relationship referencing itself or another node.
```

Question 11 of 80

Which of the following are not native security roles in Neo4j?

- A) admin
- B) publisher
- **C) developer**
- D) architect

```
Points: 0 out of 1
Correct answer: C)
```

Question 12 of 80

What parts of a Cypher query can be parameterized?

- A) Labels.
- **B) Property values.**
- **C) Map values.**
- D) Relationship types.
- **E) Index query expressions for explicit indexes.**
- **F) Literals.**

```
Points: 0 out of 2
Correct answer: B) C) E) F)
```

Question 13 of 80

What is the difference between count(*) and count(n)?

- A) There is no difference.
- **B) count(*) counts rows not values.**
- C) count(n) only counts distinct values.
- **D) count(n) does not count null values.**

```
Points: 0.5 out of 1
Correct answer: B) D)
```

Question 14 of 80

Which of the following best describes the Neo4j data model?

- A) Nodes and properties.
- **B) Nodes, relationships, and properties.**
- C) Nodes, foreign keys, and relationships.
- D) Rows, properties, and relationships.
- E) Nodes and relationships.

```
Points: 0 out of 1
Correct answer: B)
Neo4j's data model consists of nodes, relationships between the nodes, and properties on both nodes
and relationships.
```

Question 15 of 80

What prefix should be used in the connection URI when connecting an application to a Neo4j causal cluster?

- A) bolt
- B) bolt+cluster
- **C) bolt+routing**
- D) bolt+causalcluster

```
Points: 0 out of 1
Correct answer: C)
```

Question 16 of 80

Given this Cypher statement, select the answer that best describes what data is returned from the query?

```
MATCH (person:Person)
WHERE person.id = 526321
RETURN KEYS(person)
```

- **A) A collection of the property keys stored on the node with the Person label and an id property with a value of 526321.**
- B) The internal key reference pointer for the node with the Person label and an id property with a value of 526321.
- C) The index keys stored for the node with the Person label and an id property with a value of 526321.
- D) The key pointers for relationships connected to the node with the Person label and an id property with a value of 526321.

```
Points: 0 out of 1
Correct answer: A)
```

Question 17 of 80

What are the benefits of parameterized Cypher?

- **A) The query plan from previous invocations is reused, therefore it's much more performant.**
- **B) Provides protection against code injection.**
- C) Automatic query parallelization in a cluster
- D) There are no benefits. Constructing statements with literal values is as good.
- E) Parameters are only available in SQL, Cypher doesn't support it.

```
Points: 0 out of 2
Correct answer: A) B)
```

Question 18 of 80

The following Cypher statement may not return the list of all employees who work for Acme.

```
MATCH (a:Employee {id:5})-[:WORKS_FOR]->(b:Company {name:"Acme"})
RETURN a
LIMIT 10
```

Select the statement(s) below that describe why the statement may not return the list of all employees who work for Acme.

- A) We are returning only the variable a; to get the list of employee names we would have to return a.name.
- **B) The LIMIT 10 following the RETURN clause means we will only get 10 results and there may be more than 10 employees who work at Acme in the graph.**
- **C) We are matching on a node with an Employee label and with an id property of value 5, which may refer only to a single employee.**
- D) The LIST keyword needs to be used in the RETURN statement to generate the list of all nodes.

```
Points: 0 out of 1
Correct answer: B) C)
```

Question 19 of 80

Neo4j uses the Property graph model. Which of the following best describes a property graph?

- **A) Nodes and relationships define the graph while properties add context by storing relevant information in the nodes and relationships.**
- B) Property graph defines a graph meta-structure that acts as a model or schema for the data as it is entered.
- C) The Property graph is a model similar to RDF which describes how Neo4j stores resources in the database.
- D) The Property graph allows for configuration properties to define schema and structure of the graph.

```
Points: 0 out of 1
Correct answer: A)
```

Question 20 of 80

Which of the following best describes the LIMIT clause in Cypher?

- A) **LIMIT is used to limit the number rows returned from the query or passed to other parts of a query.**
- B) LIMIT is used within the WHERE clause to limit the number of relationships traversed during a query.
- C) The LIMIT clause is used when creating nodes to limit the number of relationships between two nodes.
- D) The LIMIT clause is used with a RETURN clause to limit the types of values returned from node properties.

```
Points: 0 out of 1
Correct answer: A)
```

Question 21 of 80

When using a language driver, a Session can be considered thread safe?

- A) True
- B) **False**

```
Points: 0 out of 1
Correct answer: B) Selected answer: No answer given
```

Question 22 of 80

In modeling, the use of Bi-directional relationships is a good practice when which of the following is true?

- **A) The semantics of the relationship in one direction is different from the other direction.**
- B) When you want to show the same relationship between two nodes in each direction.
- C) When a relationship between two nodes could be represented in either direction.
- D) This is never a good idea.

```
Points: 0 out of 1
Correct answer: A)
```

Question 23 of 80

The Cypher PROFILE keyword can be used for what purpose?

- **A) Entered before the statement it is used to return the query plan and execution information for a Cypher statement for performance tuning purposes.**
- B) Used when creating parameterized Cypher queries, it tells the query engine to build a query plan for later use.
- C) The PROFILE clause will detail the current statistics for the server, including node counts, relationship counts, and data size.
- D) PROFILE will identify the schema for the current database, including labels in use, relationship types, and indexes.

```
Points: 0 out of 1
Correct answer: A)
```

Question 24 of 80

When using a language driver, transactions can be executed in which access modes?

- A) ReadWriteOnce, ReadOnlyMany, or ReadWriteMany
- B) Master or Slave
- **C) Read or Write**
- D) Serializable, Read committed, or Read uncommitted

```
Points: 0 out of 1
Correct answer: C)
```

Question 25 of 80

Which of the following Cypher statements will return the number of cities in the state of California?

- **A)**

```
MATCH (:State {name:"California"})<-[:LOCATED_IN]-(city:City)
RETURN count(city)
```

- **B)**

```
MATCH (state:State)<-[:LOCATED_IN]-(city:City)
WHERE state.name="California"
RETURN count(city)
```

- C)

```
MATCH (state:State {name:"California"})
JOIN state, MATCH (city:City)
RETURN count(city)
```

- D)

```
MATCH (city:City)
FILTER relationships(LOCATED_IN)
FILTER related(:STATE {name:"California"})
RETURN count(city)
```

```
Points: 0 out of 2
Correct answer: A) B)
```

Question 26 of 80

What is the main responsibility of core servers in a causal cluster?

- A) Scaling reads
- **B) Safeguarding data**
- C) Processing analytics queries
- D) Providing the cluster topology

```
Points: 0 out of 1
Correct answer: B)
```

Question 27 of 80

Two nodes representing the same kind of thing, such as a person, can be connected to other nodes using different relationship types.

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 28 of 80

When modeling with Neo4j, which of the following best represents inferring a symmetric relationship for the below simple graph:

```
(parent:Parent)–[:PARENT_OF]–>(child:Child)
```

- **A) child is a CHILD_OF parent**
- B) parent is the CHILD_OF parent
- C) child is the PARENT_OF parent
- D) child HAS_SIBLINGS from parent

```
Points: 1 out of 1
Correct answer: A)
```

Question 29 of 80

A single node can have a relationship that points at itself?

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 30 of 80

What's the default port for the Bolt server?

- **A) 7687**
- B) 7474
- C) 7473
- D) 5000

```
Points: 1 out of 1
Correct answer: A)
```

Question 31 of 80

Which of the following best describes a Relationship in Neo4j?

- **A) A structure with a name and direction that describes the relationship between two nodes and provides structure and context to the graph.**
- B) The link between two types of nodes.
- C) A link that indicates how one type of node is, or should be connected to another type of node.
- D) A key/value pair that identifies additional nodes that a single node is related to, including direction and weight

```
Points: 0 out of 1
Correct answer: A)
```

Question 32 of 80

Neo4j requires which of the following?

- **A) Java runtime environment.**
- B) Microsoft .NET environment.
- C) A J2EE container server.
- D) A SAN storage system.
- E) None of the above.

```
Points: 0 out of 1
Correct answer: A)
Neo4j requires a Java runtime environment for deployment.
```

Question 33 of 80

Foreign keys are necessary in a graph database. These allow you to determine which nodes are related.

- A) True
- **B) False**

```
Points: 0 out of 1
Correct answer: B)
Nodes are connected by relationships. Foreign keys are not needed.
```

Question 34 of 80

Is ORDER BY a valid Cypher clause?

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 35 of 80

In modeling, the concept of an intermediate node is used in what situation?

- **A) To connect more than two nodes in a single context, such as where a person worked and in what role during what timeframe.**
- B) As a sub-node to represent complex entities as multiple nodes.
- C) In support of linked lists to identify next relationships.
- D) Before or after Primary nodes for quick traversals through the graph.

```
Points: 0 out of 1
Correct answer: A)
```

Question 36 of 80

Which function allows you to create a list of values as result of an aggregation?

- **A) The collect() function.**
- B) The toList() function.
- C) The aggregate() function.
- D) The values() function.

```
Points: 0 out of 1
Correct answer: A)
```

Question 37 of 80

What's the correct exception to throw to signal an error from within a procedure?

- **A) RuntimeException**
- B) ProcedureException
- C) NotFoundException
- D) TransientTransactionFailureException

```
Points: 0 out of 1
Correct answer: A)
```

Question 38 of 80

Which of the following best describes what the below Cypher statement will do?

```
MATCH (city:City {name: "San Mateo"})
MERGE (state:State {name: "California"})
MERGE (city)-[:LOCATED_IN]->(state)
RETURN city, state
```

- A) **If there is a City node with name "San Mateo", uniquely create a :LOCATED_IN relationship to a State node with name "California," creating the State node if it does not already exist.**
- B) The :LOCATED_IN relationship is only created if there is both a City node with name "San Mateo" and a State node with name "California."
- C) The :LOCATED_IN relationship is overwritten only if there is both a City node with name "San Mateo" and a State node with name "California."
- D) The :LOCATED_IN relationship is matched only if there is both a City node with name "San Mateo" and a State node with name "California."

```
Points: 0 out of 1
Correct answer: A)
```

Question 39 of 80

Online backups allow for backing up a running instance of Neo4j, which is only available with the Enterprise Edition.

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 40 of 80

Which procedure can be run to get a list of all procedures in the DBMS?

- A) db.procedures()
- B) dbms.showProcedures()
- C) db.listProcedures()
- **D) dbms.procedures()**

```
Points: 0 out of 1
Correct answer: D)
```

Question 41 of 80

Which keyword in the RETURN clause will return only one instance of each item in a result set?

- A) UNIQUE
- **B) DISTINCT**
- C) SINGLE
- D) FIRST

```
Points: 0 out of 1
Correct answer: B)
```

Question 42 of 80

The four building blocks of a Neo4j Graph Database are:

1. Nodes
2. Relationships
3. Properties
4. Labels

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 43 of 80

Which of the following schema objects does Neo4j include?

- A) Tables.
- B) Columns.
- C) Materialized views.
- D) Rows.
- **E) None of the above.**

```
Points: 0 out of 1
Correct answer: E)
Neo4j is a graph database and does not enforce relational schema objects.
```

Question 44 of 80

In a Neo4j database, which of the below best describes what Nodes are used for?

- **A) Used to represent entities and complex value types in the graph.**
- B) As a table structure that identifies like entities and groups them together.
- C) As a reference holder for keys, which also store values, and foreign key links to other nodes.
- D) As endpoints on either side of a relationship, used to bring together relationships in the graph.

```
Points: 0 out of 1
```

Correct answer: A)

Question 45 of 80

In Neo4j modeling, a timeline tree is a recommended approach for representing time and connecting discrete events with no natural relationship to other events where you need to find events by granularity of time.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 46 of 80

Neo4j allows for undirected relationships between nodes.

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)
Neo4j requires each relationship to have a direction and type.

Question 47 of 80

Which of the following Cypher statements would return the total population in all cities located in California?

- **A)**

```
MATCH (:STATE {name:"California"})<-[:LOCATED_IN]-(city:CITY)
RETURN sum(city.population)
```

- B)

```
MATCH (city:CITY)
sum(city.population) as total
WHERE (city)-[:LOCATED_IN]->(:STATE {name="California"})
RETURN total
```

- C)

```
SUM (:CITY.population)
WHERE city.relationships(:SATE.name="California")
RETURN
```

- D)

```
MATCH (state:STATE {name:"California"})
MATCH (city:CITY)
```

```
JOIN state,city
RETURN SUM(city.population)
```

```
Points: 0 out of 1
Correct answer: A)
```

Question 48 of 80

Which settings determine the maximum memory use of Neo4j?

- **A) dbms.pagecache.memory and wrapper.java.maxmemory.**
- B) object_cache_size, file_cache_size, heap_size.
- C) dbms.pagecache.size
- D) java.io.heapsize, neo4j.db.cachesize
- E) Java option -Xmx

```
Points: 0 out of 1
Correct answer: A)
```

Question 49 of 80

Which statements regarding page cache are true:

- **A) The page cache uses off-heap memory.**
- B) The Java heap contains the page cache.
- **C) If possible you should assign enough RAM to "dbms.pagecache.memory" to hold the full graph.**
- D) The page cache is implemented in C++ for maximum throughput.
- **E) The page cache has a minimal of overhead compared to the binary storage on disc.**

```
Points: 0 out of 2
Correct answer: A) C) E)
```

Question 50 of 80

Neo4j is written in C++.

- A) True
- **B) False**

```
Points: 0 out of 1
Correct answer: B)
Neo4j is written in Java.
```

Question 51 of 80

Which of the following statements best describes properties?

- **A) Properties are the key-value pairs on both nodes and relationships.**

- B) Properties are the key-value pairs on nodes only.
- C) Properties are the key-value pairs on relationships only.
- D) None of the above.

```
Points: 0 out of 1
Correct answer: A)
Properties are the key-value pairs on both nodes and relationships.
```

Question 52 of 80

Select the Cypher statements below that will find all of Jason's friends and set their verified property to true?

- **A)**

```
MATCH (a:Person {name:"Jason"})-[:FRIEND]-(b:Person)
WITH COLLECT(b) AS friends
FOREACH (n IN friends | SET n.verified = TRUE)
```

- B)

```
MATCH (a:Person)-[:FRIEND]->(b)
UPDATE b
SET b.verified = TRUE
```

- C)

```
MATCH (a:Person)-[:FRIEND]->(b)
WHERE a.name = "Jason"
UPDATE b
SET b.verified = TRUE
```

- **D)**

```
MATCH (a:Person {name:"Jason"})-[:FRIEND]-(b:Person)
SET b.verified = TRUE
```

```
You did not select all available correct options.
Points: 0 out of 1
Correct answer: A) D)
```

Question 53 of 80

Which export formats does the Neo4j Browser support?

- **A) PNG**
- B) XML
- C) GraphML
- D) XLS

- **E) CSV**
- **F) SVG**

```
Points: 0 out of 1
Correct answer: A) E) F)
```

Question 54 of 80

Which of the following best describes the options available for loading data into Neo4j?

- **A) Neo4j's neo4j-import tool, Cypher LOAD CSV clause, or batch operations against the REST endpoint.**
- B) Commercial ETL tools that can transform and load data in graph format.
- C) Direct database connections from other database tools to load data in directly.
- D) The Cypher BULK LOAD clause to import data from an existing text file.

```
Points: 0 out of 1
Correct answer: A)
```

Question 55 of 80

Which of the following authentication providers does Neo4j have support for?

- A) Google
- **B) Kerberos**
- **C) LDAP**
- D) LinkedIn

```
Points: 0 out of 1
Correct answer: B) C)
```

Question 56 of 80

Given the below MATCH clause, which of the following is the correct way to return the relationship type of all relationships bound to the variable b?

```
MATCH (a)—[b]—>(c)
```

- A) RETURN b.type
- B) RETURN b
- **C) RETURN TYPE(b)**
- D) None of these are correct.

```
Points: 0 out of 1
Correct answer: C)
```

Question 57 of 80

Adding a label to a node will automatically index the name property.

- A) True
- **B) False**

```
Points: 0 out of 1
Correct answer: B)
Automatic indexing of properties only occurs when the configuration is enabled or the index or
uniqueness constraint is applied.
```

Question 58 of 80

How do you define ";" as field terminator in LOAD CSV?

- **A) LOAD CSV FROM "url" AS row FIELDTERMINATOR ";"**
- B) LOAD CSV FROM "url" AS row TERMINATED BY ";"
- C) LOAD CSV DELIMETER ";" FROM "url" AS row
- D) LOAD CSV FROM "url" AS row WITH split(row, ";") as fields

```
Points: 0 out of 1
Correct answer: A)
```

Question 59 of 80

Which of the following Cypher statements will list all the labels in the database?

- A) RETURN db.labels()
- **B) CALL db.labels()**
- C) RETURN CALL db.labels()
- **D) CALL db.labels() YIELD label RETURN label**

```
Points: 0 out of 1
Correct answer: B) D)
```

Question 60 of 80

Given a model consisting of nodes with the Person label connected by relationships with type KNOWS, select the statement below that will match on both Sarah's friends and her friends of friends?

- A)

```
MATCH (sarah:Person)-[:KNOWS]->(friend:Person)
WHERE sarah.name = "Sarah"
RETURN friend
```

- **B)**

```
MATCH (sarah:Person)-[:KNOWS*1..2]->(friend:Person)
WHERE sarah.name = "Sarah"
RETURN friend
```

- **C)**

```
MATCH (sarah:Person)—[:KNOWS]—>(friend:Person)
OPTIONAL MATCH (friend)—[:KNOWS]—>(fof:Person)
WHERE sarah.name = "Sarah"
RETURN friend, fof
```

- D)

```
All of these are correct.
```

```
Points: 0 out of 2
Correct answer: B) C)
```

Question 61 of 80

Referential integrity is maintained in Neo4j because you cannot delete a node that has existing relationships attached.

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
Nodes cannot be deleted if they still have relationships attached to them.
```

Question 62 of 80

In the following MATCH clause, which of the elements represents the relationship between two nodes?

```
MATCH (a)—[b]—>(c)
```

- A) (a)
- B) (a)-
- **C) -[b]->**
- D) (c)
- E) ->(c)

```
Points: 0 out of 1
Correct answer: C)
```

Question 63 of 80

Of the following, which are reasons why Labels are used?

- **A) Used to represent entities, such as users, products, or company.**

- B) **Used to group like nodes together.**
- C) Used to represent tables in the database.

* **D) Used to associate indexes and constraints with groups of nodes.**

```
Points: 0 out of 3
Correct answer: A) B) D)
```

Question 64 of 80

In Neo4j modeling, the use of a Linked List is helpful for which of the following reasons?

- A) **To link entities together in a sequence, such as a series of events.**
- B) **Identifying the start and end nodes of a series of events.**
- C) To chain together nodes with similar labels.
- D) To always start a graph traversal at a certain point.

```
Points: 0 out of 2
Correct answer: A) B)
```

Question 65 of 80

When running an Online backup, Neo4j will automatically select to run a full or incremental backup based on whether an existing backup is present at the specified backup location and there are enough log files since the last backup to support a incremental backup?

- A) **True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
```

Question 66 of 80

Property values can be the following:

- A) **Numbers**
- B) **boolean values**
- C) **byte[]**
- D) **Strings**
- E) **Arrays of Strings**
- F) List<String>
- G) **Date**
- H) Nested Documents

```
Points: 0 out of 1
Correct answer: A) B) C) D) E) G)
```

Question 67 of 80

Neo4j can be deployed...

- A) ...embedded in a Java application.
- B) ...as a standalone server.
- **C) ...embedded in a Java application or as a standalone server.**
- D) ...with an app server only.
- E) None of the above.

```
Points: 0 out of 1
Correct answer: C)
Neo4j can either be embedded in a Java application or deployed as a standalone server.
```

Question 68 of 80

Which of the following Raft roles are not played by core servers in a causal cluster?

- A) Leader
- B) Follower
- C) **Master**
- D) Candidate
- E) **Slave**

```
Points: 0 out of 1
Correct answer: C) E)
```

Question 69 of 80

How can we ensure that the same relationship cannot be created twice between two nodes?

- A) Use the DISTINCT keyword when creating the relationship.
- B) Use the UNIQUE keyword when creating the relationship.
- C) Use the ONLY keyword when creating the relationship.
- **D) Use the MERGE keyword when creating the relationship.**

```
Points: 0 out of 2
Correct answer: D)
```

Question 70 of 80

What type of database is Neo4j?

- A) Key-value store.
- B) Document database.
- **C) Graph database.**
- D) Relational database.
- E) Semantic database.

```
Points: 0 out of 1
Correct answer: C)
Neo4j is a graph database.
```

Question 71 of 80

To transform a collection of nodes or relationships into a collection of properties or results of a function or expression, which Cypher clause would you use?

- **A) EXTRACT**
- B) GET
- C) REDUCE
- D) FILTER
- E) KEYS
- F) ALL
- G) MERGE

```
Points: 0 out of 1
Correct answer: A)
```

Question 72 of 80

All nodes with the same label must have the same property keys.

- A) True
- B) **False**

```
Points: 0 out of 1
Correct answer: B)
Neo4j does not enforce by default that all nodes with the same label must have the same property keys.
This is only guaranteed with Property Existence Constraints.
```

Question 73 of 80

On a Neo4j instance participating in a cluster, which of the following configuration settings is used to define the list of other known instances that cluster?

- **A) causal_clustering.initial_discovery_members**
- B) causal_clustering.host_list
- C) causal_clustering.cluster_members
- D) causal_clustering.cluster_instances

```
Points: 0 out of 1
Correct answer: A)
```

Question 74 of 80

Which configuration property must be set to allow a procedure access to unsupported restricted APIs in the database?

- A) dbms.security.unrestricted_access
- **B) dbms.security.procedures.unrestricted**
- C) dbms.allow_all_operations
- D) dbms.procedures.access_all

```
Points: 0 out of 1
Correct answer: B)
```

Question 75 of 80

In a causal cluster, a read replica can process write queries

- A) True
- B) **False**

```
Points: 0 out of 1
Correct answer: B)
```

Question 76 of 80

When using a language driver, which forms of transaction can be automatically replayed on failure?

- A) Auto-commit transactions
- **B) Transaction functions**
- C) Explicit transactions

```
Points: 0 out of 1
Correct answer: B)
```

Question 77 of 80

Which of the following best describes a Neo4j incremental backup?

- A) A snapshot of the transaction log from the last backup point is copied to the backup location so it can be replayed during a restore.
- **B) An incremental backup is performed whenever an existing backup directory is specified. The backup tool will then copy any new transactions from the Neo4j server and apply them to the backup. The result will be an updated backup that is consistent with the current server state.**
- C) An incremental backup is a scheduled backup process in Neo4j that copies transactions from the database to a backup location for future recovery purposes.
- D) Neo4j does not have an incremental backup capability.

```
Points: 0 out of 1
Correct answer: B)
```

Question 78 of 80

Neo4j is available with both open-source and commercial licenses.

- **A) True**
- B) False

```
Points: 0 out of 1
Correct answer: A)
Neo4j can be used with both the open-source license or a commercial license which includes technical
support.
```

Question 79 of 80

Which of the following best describes Cypher, Neo4j's graph query language?

- A) It's a SQL plugin for Neo4j.
- B) It is a regular expression-like programming language for interfacing with Neo4j.
- **C) It is a declarative query language designed for graph pattern matching and traversals.**
- D) It is a procedural programming language for interfacing with Neo4j.

```
Points: 1 out of 1
Correct answer: C)
Cypher is a declarative query language designed for graph pattern matching and traversals.
```

Question 80 of 80

What information can we not find out by executing the dbms.listQueries() procedure?

- A) Bytes allocated for the executing query
- B) Planner used by the query
- C) Username of the user who is executing the query.
- D) **IP address of the machine that is executing the query**

```
Points: 0 out of 1
Correct answer: D)
```

Retrieved from 'https://wiki.glitchdata.com/index.php?title=Neo4J:_Sample_Questions_3&oldid=23144'