

Neo4J: Sample Questions 5

Contents

Question 1 of 80
Question 2 of 80
Question 3 of 80
Question 4 of 80
Question 5 of 80
Question 6 of 80
Question 7 of 80
Question 8 of 80
Question 9 of 80
Question 10 of 80
Question 11 of 80
Question 12 of 80
Question 13 of 80
Question 14 of 80
Question 15 of 80
Question 16 of 80
Question 17 of 80
Question 18 of 80
Question 19 of 80
Question 20 of 80
Question 21 of 80
Question 22 of 80
Question 23 of 80
Question 24 of 80
Question 25 of 80
Question 26 of 80
Question 27 of 80
Question 28 of 80
Question 29 of 80
Question 30 of 80
Question 31 of 80

Question 32 of 80

Question 33 of 80

Question 34 of 80

Question 35 of 80

Question 36 of 80

Question 37 of 80

Question 38 of 80

Question 39 of 80

Question 40 of 80

Question 41 of 80

Question 42 of 80

Question 43 of 80

Question 44 of 80

Question 45 of 80

Question 46 of 80

Question 47 of 80

Question 48 of 80

Question 49 of 80

Question 50 of 80

Question 51 of 80

Question 52 of 80

Question 54 of 80

Question 55 of 80

Question 56 of 80

Question 57 of 80

Question 58 of 80

Question 59 of 80

Question 60 of 80

Question 61 of 80

Question 62 of 80

Question 63 of 80

Question 64 of 80

Question 65 of 80

Question 66 of 80

Question 67 of 80

Question 68 of 80

Question 69 of 80

Question 70 of 80

Question 71 of 80

Question 72 of 80

Question 73 of 80

Question 74 of 80

Question 75 of 80

Question 76 of 80

Question 77 of 80

Question 78 of 80

Question 79 of 80

Question 80 of 80

Question 1 of 80

Which of the following best describes the options available for loading data into Neo4j?

- **A) Neo4j's neo4j-import tool, Cypher LOAD CSV clause, or batch operations against the REST endpoint.**
- B) Commercial ETL tools that can transform and load data in graph format.
- C) Direct database connections from other database tools to load data in directly.
- D) The Cypher BULK LOAD clause to import data from an existing text file.

Points: 0 out of 1
Correct answer: A)

Question 2 of 80

You can create the same type of relationship between two nodes with each relationship having different properties?

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 3 of 80

dbms.mode is a configuration parameter that allows you to set whether the instance runs in SINGLE, CORE, or READ_REPLICA mode.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 4 of 80

Which of the following best describes what the below Cypher statement will do?

```
MATCH (city:City {name: "San Mateo"})
MERGE (state:State {name: "California"})
MERGE (city)-[:LOCATED_IN]->(state)
RETURN city, state
```

- **A) If there is a City node with name "San Mateo", uniquely create a :LOCATED_IN relationship to a State node with name "California," creating the State node if it does not already exist.**
- B) The :LOCATED_IN relationship is only created if there is both a City node with name "San Mateo" and a State node with name "California."
- C) The :LOCATED_IN relationship is overwritten only if there is both a City node with name "San Mateo" and a State node with name "California."
- D) The :LOCATED_IN relationship is matched only if there is both a City node with name "San Mateo" and a State node with name "California."

Points: 0 out of 1
Correct answer: A)

Question 5 of 80

Which of the following best describes the LIMIT clause in Cypher?

- **A) LIMIT is used to limit the number rows returned from the query or passed to other parts of a query.**
- B) LIMIT is used within the WHERE clause to limit the number of relationships traversed during a query.
- C) The LIMIT clause is used when creating nodes to limit the number of relationships between two nodes.
- D) The LIMIT clause is used with a RETURN clause to limit the types of values returned from node properties.

Points: 0 out of 1
Correct answer: A)

Question 6 of 80

What's the default port for the Bolt server?

- **A) 7687**
- B) 7474
- C) 7473
- D) 5000

Points: 1 out of 1
Correct answer: A)

Question 7 of 80

Which of the following are possible ways to monitor the state of a Neo4j instance?

- **A) Using the metrics plugin available with Enterprise Edition.**
- **B) Using the :sysinfo command in Neo4j Browser to view information about the current server instance.**

- C) Using the Neo4j monitoring-shell service to remotely run monitoring commands against Neo4j instances.
- D) Using Neo4j's optionally available, web-based Manager application for full remote monitoring and management capabilities.

Points: 0 out of 2
Correct answer: A) B)

Question 8 of 80

Which of the following best describes the CONNECT_BY clause in Cypher

- A) The CONNECT_BY clause is used to join nodes when they are connected by varying relationship depths.
- B) The CONNECT_BY clause is used when constructing a graph tree structure where you want to define the nodes that a leaf node is connected to.
- C) The CONNECT_BY clause is used in Cypher to limit the pattern to only certain relationship types.
- D) The CONNECT_BY clause is a constraint that ensures only certain relationships can connect two nodes with specific labels together.
- **E) CONNECT_BY is not a valid Cypher clause.**

Points: 0 out of 1
Correct answer: E)

Question 9 of 80

What, specifically, does the following expression do?: size((n)-[:TYPE]->()) ?

- **A) It counts the number of relationships with this type and direction?**
- B) This is not valid Cypher.
- **C) It reads the degree-value for this relationship-type and direction directly from the node record.**
- D) It executes a MATCH first and then returns the size of the path - collection.

Points: 0 out of 2
Correct answer: A) C)

Question 10 of 80

Which configuration property must be set to allow a procedure access to unsupported restricted APIs in the database?

- A) dbms.security.unrestricted_access
- **B) dbms.security.procedures.unrestricted**
- C) dbms.allow_all_operations
- D) dbms.procedures.access_all

Points: 0 out of 1
Correct answer: B)

Question 11 of 80

In a causal cluster, a read replica can process write queries

- A) True
- B) **False**

Points: 0 out of 1
Correct answer: B)

Question 12 of 80

Which of the following would be the correct Cypher syntax to create an index on the id property for nodes with the Customer label?

- A) **CREATE INDEX ON :Customer(id);**
- B) CREATE INDEX ON :Customer.id;
- C) CREATE INDEX ON Customer(id);
- D) CREATE GRAPH INDEX ON :Customer.id;

Points: 0 out of 1
Correct answer: A)

Question 13 of 80

When modeling relationships between two nodes, you should use a NULL relationship to represent the absence of a relationship between two nodes.

- A) True
- B) **False**

Points: 0 out of 1
Correct answer: B)
The absence of a relationship between two nodes is sufficient to represent no connection between them.

Question 14 of 80

Which of the following best describes the Neo4j data model?

- A) Nodes and properties.
- B) **Nodes, relationships, and properties.**
- C) Nodes, foreign keys, and relationships.
- D) Rows, properties, and relationships.
- E) Nodes and relationships.

Points: 0 out of 1
Correct answer: B)
Neo4j's data model consists of nodes, relationships between the nodes, and properties on both nodes and relationships.

Question 15 of 80

Which of the following best describes Cypher, Neo4j's graph query language?

- A) It's a SQL plugin for Neo4j.
- B) It is a regular expression-like programming language for interfacing with Neo4j.
- **C) It is a declarative query language designed for graph pattern matching and traversals.**
- D) It is a procedural programming language for interfacing with Neo4j.

Points: 1 out of 1

Correct answer: C)

Cypher is a declarative query language designed for graph pattern matching and traversals.

Question 16 of 80

Relationships are defined with regard to node instances, not classes of nodes.

- **A) True**
- B) False

Points: 0 out of 1

Correct answer: A)

Question 17 of 80

Which of the following Cypher statements will return the number of cities in the state of California?

- **A)**

```
MATCH (:State {name:"California"})<-[:LOCATED_IN]-(city:City)
RETURN count(city)
```

- **B)**

```
MATCH (state:State)<-[:LOCATED_IN]-(city:City)
WHERE state.name="California"
RETURN count(city)
```

- **C)**

```
MATCH (state:State {name:"California"})
JOIN state, MATCH (city:City)
RETURN count(city)
```

- **D)**

```
MATCH (city:City)
```

```

FILTER relationships(LOCATED_IN)
FILTER related(:STATE {name:"California"})
RETURN count(city)

```

Points: 0 out of 2
Correct answer: A) B)

Question 18 of 80

Neo4j is written in C++.

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)
Neo4j is written in Java.

Question 19 of 80

Which function allows you to create a list of values as result of an aggregation?

- **A) The collect() function.**
- B) The toList() function.
- C) The aggregate() function.
- D) The values() function.

Points: 0 out of 1
Correct answer: A)

Question 20 of 80

Is ORDER BY a valid Cypher clause?

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 21 of 80

In which directory would you find the neo4j.conf file that defines the configuration of the server?

- **A) "conf" directory**
- B) "properties" directory

- C) "bin" directory
- D) "data" directory

Points: 0 out of 1
Correct answer: A)

Question 22 of 80

To transform a collection of nodes or relationships into a collection of properties or results of a function or expression, which Cypher clause would you use?

- **A) EXTRACT**
- B) GET
- C) REDUCE
- D) FILTER
- E) KEYS
- F) ALL
- G) MERGE

Points: 0 out of 1
Correct answer: A)

Question 23 of 80

How do you perform an aggregation in Cypher?

- A) Using the GROUP BY keyword.
- **B) Using at least one aggregation function.**
- C) Defining grouping keys with WITH.
- D) With the AGGREGATE keyword.

Points: 0 out of 1
Correct answer: B)

Question 24 of 80

The maximum number of instances in a Neo4j cluster is 9.

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)
Incorrect, there is no hard limit on the number of instances in a cluster.

Question 25 of 80

Select the Cypher statements below that will find all of Jason's friends and set their verified property to true?

▪ A)

```
MATCH (a:Person {name:"Jason"})-[:FRIEND]-(b:Person)
WITH COLLECT(b) AS friends
FOREACH (n IN friends | SET n.verified = TRUE)
```

▪ B)

```
MATCH (a:Person)-[:FRIEND]->(b)
UPDATE b
SET b.verified = TRUE
```

▪ C)

```
MATCH (a:Person)-[:FRIEND]->(b)
WHERE a.name = "Jason"
UPDATE b
SET b.verified = TRUE
```

▪ D)

```
MATCH (a:Person {name:"Jason"})-[:FRIEND]-(b:Person)
SET b.verified = TRUE
```

You did not select all available correct options.
Points: 0 out of 1
Correct answer: A) D)

Question 26 of 80

A single node can have a relationship that points at itself?

- A) True
- B) False

Points: 0 out of 1
Correct answer: A)

Question 27 of 80

Which of the following are not native security roles in Neo4j?

- A) admin
- B) publisher
- C) developer
- D) architect

Points: 0 out of 1
Correct answer: C)

Question 28 of 80

In the following MATCH clause, which of the elements represents the relationship between two nodes?

MATCH (a)-[b]->(c)

- A) (a)
- B) (a)-
- **C) -[b]->**
- D) (c)
- E) ->(c)

Points: 0 out of 1
Correct answer: C)

Question 29 of 80

Which keyword in the RETURN clause will return only one instance of each item in a result set?

- A) UNIQUE
- **B) DISTINCT**
- C) SINGLE
- D) FIRST

Points: 0 out of 1
Correct answer: B)

Question 30 of 80

Neo4j can be deployed...

- A) ...embedded in a Java application.
- B) ...as a standalone server.
- **C) ...embedded in a Java application or as a standalone server.**
- D) ...with an app server only.
- E) None of the above.

Points: 0 out of 1
Correct answer: C)

Neo4j can either be embedded in a Java application or deployed as a standalone server.

Question 31 of 80

To restore a Neo4j instance from a backup, which of the following is accurate?

- **A) The Neo4j backups are fully functional databases. To use a backup, all you need to do replace your database folder with the backup. Just make sure the database isn't running while replacing the folder.**
- B) The Neo4j "backup-restore" command will copy the compressed backup archive files and logical logs to the existing instance database folder, replay the transactions, and bring the database back online.
- C) Issuing a "restore" command with a timestamp will recover the Neo4j instance up to that particular point in time, before any transactions corrupted the database files.
- D) Shut down the current instance, replace the database folder with the backup, and then startup the Neo4j instance in recovery mode so the database is fully recovered.

Points: 0 out of 1
Correct answer: A)

Question 32 of 80

What is the difference between count(*) and count(n)?

- A) There is no difference.
- **B) count(*) counts rows not values.**
- C) count(n) only counts distinct values.
- **D) count(n) does not count null values.**

Points: 0.5 out of 1
Correct answer: B) D)

Question 33 of 80

The four building blocks of a Neo4j Graph Database are:

1. Nodes
2. Relationships
3. Properties
4. Labels

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 34 of 80

Given a model consisting of nodes with the Person label connected by relationships with type KNOWS, select the statement below that will match on both Sarah's friends and her friends of friends?

- A)

```
MATCH (sarah:Person)-[:KNOWS]->(friend:Person)
WHERE sarah.name = "Sarah"
RETURN friend
```

▪ B)

```
MATCH (sarah:Person)-[:KNOWS*1..2]->(friend:Person)
WHERE sarah.name = "Sarah"
RETURN friend
```

▪ C)

```
MATCH (sarah:Person)-[:KNOWS]->(friend:Person)
OPTIONAL MATCH (friend)-[:KNOWS]->(fof:Person)
WHERE sarah.name = "Sarah"
RETURN friend, fof
```

▪ D)

All of these are correct.

Points: 0 out of 2
Correct answer: B) C)

Question 35 of 80

Select the Cypher statements below that will delete a node with an id of 3563 and all of its possibly connected relationships?

▪ A)

```
MATCH (a:Thing {id:3563})
OPTIONAL MATCH (a)-[r]-()
DELETE a, r
```

▪ B)

```
MATCH (a:Thing {id:3563})
DELETE a
```

▪ C)

```
MATCH (a:Thing {id:3563})
DETACH DELETE a
```

▪ D)

```
MATCH (a:Thing {id:3563})-[r]-(b)
DELETE a, r
```

Points: 0 out of 1
Correct answer: A) C)

Question 36 of 80

Which of the following statements best describes properties?

- **A) Properties are the key-value pairs on both nodes and relationships.**
- B) Properties are the key-value pairs on nodes only.
- C) Properties are the key-value pairs on relationships only.
- D) None of the above.

Points: 0 out of 1
Correct answer: A)
Properties are the key-value pairs on both nodes and relationships.

Question 37 of 80

Referential integrity is maintained in Neo4j because you cannot delete a node that has existing relationships attached.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)
Nodes cannot be deleted if they still have relationships attached to them.

Question 38 of 80

In modeling, the concept of an intermediate node is used in what situation?

- **A) To connect more than two nodes in a single context, such as where a person worked and in what role during what timeframe.**
- B) As a sub-node to represent complex entities as multiple nodes.
- C) In support of linked lists to identify next relationships.
- D) Before or after Primary nodes for quick traversals through the graph.

Points: 0 out of 1
Correct answer: A)

Question 39 of 80

Which of the following Cypher statements will find all procedures with the "db." prefix?

- **A)**

```
CALL dbms.procedures() yield name, signature, description
WHERE name starts with "db."
return name, signature, description
```

- B)

```
CALL dbms.procedures() yield *  
WHERE name starts with "db."  
return name, signature, description
```

- C)

```
CALL dbms.procedures("db.*")  
return name, signature, description
```

- D)

```
CALL db.procedures() yield name, signature, description  
WHERE name starts with "db."  
return name, signature, description
```

Points: 0 out of 1
Correct answer: A)

Question 40 of 80

Property values can be the following:

- A) Numbers
- B) boolean values
- C) byte[]
- D) Strings
- E) Arrays of Strings
- F) List<String>
- G) Date
- H) Nested Documents

Points: 0 out of 1
Correct answer: A) B) C) D) E) G)

Question 41 of 80

Which of the following best describes a Relationship in Neo4j?

- A) A structure with a name and direction that describes the relationship between two nodes and provides structure and context to the graph.
- B) The link between two types of nodes.
- C) A link that indicates how one type of node is, or should be connected to another type of node.
- D) A key/value pair that identifies additional nodes that a single node is related to, including direction and weight

Points: 0 out of 1
Correct answer: A)

Question 42 of 80

It is possible for a node to have more than one label.

- A) True
- B) False

Points: 0 out of 1
Correct answer: A)
Nodes can have 0 or more labels.

Question 43 of 80

Two nodes representing the same kind of thing, such as a person, can be connected to other nodes using different relationship types.

- A) True
- B) False

Points: 0 out of 1
Correct answer: A)

Question 44 of 80

The following Cypher statement may not return the list of all employees who work for Acme.

```
MATCH (a:Employee {id:5})-[:WORKS_FOR]->(b:Company {name:"Acme"})
RETURN a
LIMIT 10
```

Select the statement(s) below that describe why the statement may not return the list of all employees who work for Acme.

- A) We are returning only the variable a; to get the list of employee names we would have to return a.name.
- **B) The LIMIT 10 following the RETURN clause means we will only get 10 results and there may be more than 10 employees who work at Acme in the graph.**
- **C) We are matching on a node with an Employee label and with an id property of value 5, which may refer only to a single employee.**
- D) The LIST keyword needs to be used in the RETURN statement to generate the list of all nodes.

Points: 0 out of 1
Correct answer: B) C)

Question 45 of 80

Which of the following best describes the OPTIONAL MATCH clause in Cypher?

- A) OPTIONAL MATCH is not a Cypher clause.
- **B) The OPTIONAL MATCH searches for a described pattern that may or may not exist, assigning NULL to**

any identifiers in the pattern that do not exist.

- C) The OPTIONAL MATCH clause will take a set of property values and optionally match them against all nodes in the database.
- D) OPTIONAL MATCH provides parameter placeholders for Cypher queries. It holds a parameterized query and then optionally matches it against the graph with the values supplied by the client.

Points: 0 out of 1
Correct answer: B)

Question 46 of 80

What's the correct exception to throw to signal an error from within a procedure?

- **A) RuntimeException**
- B) ProcedureException
- C) NotFoundException
- D) TransientTransactionFailureException

Points: 0 out of 1
Correct answer: A)

Question 47 of 80

In Neo4j modeling, the use of a Linked List is helpful for which of the following reasons?

Correct answer: A) B) Selected answer: No answer given A) To link entities together in a sequence, such as a series of events. B) Identifying the start and end nodes of a series of events. C) To chain together nodes with similar labels. D) To always start a graph traversal at a certain point.

You did not select all available correct options. Points: 0 out of 2

Question 48 of 80

Labels are best described as:

Correct answer: C) Selected answer: No answer given A) Unique tags on each node for fast lookups. B) Special types of node or relationship properties. C) Tags that are used to group nodes into sets. D) Table names in the graph database.

Points: 0 out of 1

Feedback Incorrect. Labels are a tags that are used to group nodes into sets.

Question 49 of 80

Neo4j uses the Property graph model. Which of the following best describes a property graph?

- **A) Nodes and relationships define the graph while properties add context by storing relevant information in the nodes and relationships.**
- B) Property graph defines a graph meta-structure that acts as a model or schema for the data as it is entered.
- C) The Property graph is a model similar to RDF which describes how Neo4j stores resources in the database.
- D) The Property graph allows for configuration properties to define schema and structure of the graph.

Points: 0 out of 1
Correct answer: A)

Question 50 of 80

Which of the following Cypher statements would return the total population in all cities located in California?

- **A)**

```
MATCH (:STATE {name:"California"})<-[:LOCATED_IN]-(city:CITY)
RETURN sum(city.population)
```

- B)

```
MATCH (city:CITY)
sum(city.population) as total
WHERE (city)-[:LOCATED_IN]->(:STATE {name="California"})
RETURN total
```

- C)

```
SUM (:CITY.population)
WHERE city.relationships(:STATE.name="California")
RETURN
```

- D)

```
MATCH (state:STATE {name:"California"})
MATCH (city:CITY)
JOIN state,city
RETURN SUM(city.population)
```

Points: 0 out of 1
Correct answer: A)

Question 51 of 80

We can add properties to an existing node by using which Cypher keyword?

- A) ADD
- B) UPDATE
- **C) SET**
- D) CREATE

Points: 0 out of 1
Correct answer: C)

Question 52 of 80

What things can you do in a user defined procedure that you can't do in a user defined function?

Correct answer: B) C) Selected answer: No answer given A) Execute Cypher queries B) Write to the database C) Return multiple values D) Count the number of nodes in the database

You did not select all available correct options. Points: 0 out of 1

What prefix should be used in the connection URI when connecting an application to a Neo4j causal cluster?

- A) bolt
- B) bolt+cluster
- **C) bolt+routing**
- D) bolt+causalcluster

Points: 0 out of 1
Correct answer: C)

Question 54 of 80

The only Neo4j instances that can join a cluster are those that are initially configured on the Leader to be allowed to join.

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)

Question 55 of 80

Given this Cypher statement, select the answer that best describes what data is returned from the query?

```
MATCH (person:Person)
WHERE person.id = 526321
RETURN KEYS(person)
```

- **A) A collection of the property keys stored on the node with the Person label and an id property with a value of 526321.**
- B) The internal key reference pointer for the node with the Person label and an id property with a value of 526321.

- C) The index keys stored for the node with the Person label and an id property with a value of 526321.
- D) The key pointers for relationships connected to the node with the Person label and an id property with a value of 526321.

Points: 0 out of 1
Correct answer: A)

Question 56 of 80

An Unmanaged Extension is best described by which of the following?

- **A) Unmanaged Extensions provide finer grained control over your application's interactions with Neo4j than Cypher by allowing you to write server-side code using Neo4j's Java API's and access the extension through REST calls.**
- B) Unmanaged Extensions are references to the HTTP REST API's that currently do not require authentication for reading or writing data to the graph.
- C) Unmanaged Extensions are custom application services that bypass Neo4j's Cypher interface and directly interact with data stored on disk.
- D) Unmanaged Extensions are custom Cypher extensions that provide end-user specific functional capabilities to Cypher.

Points: 0 out of 1
Correct answer: A)

Question 57 of 80

Which of the following languages have an officially supported Bolt driver?

- **A) Java**
- B) Go
- C) R
- D) PHP
- **E) Python**

Points: 0 out of 1
Correct answer: A) E)

Question 58 of 80

Of the following, which are reasons why Labels are used?

- **A) Used to represent entities, such as users, products, or company.**
- **B) Used to group like nodes together.**
- C) Used to represent tables in the database.
- * **D) Used to associate indexes and constraints with groups of nodes.**

Points: 0 out of 3
Correct answer: A) B) D)

Question 59 of 80

In modeling, the use of Bi-directional relationships is a good practice when which of the following is true?

- **A) The semantics of the relationship in one direction is different from the other direction.**
- B) When you want to show the same relationship between two nodes in each direction.
- C) When a relationship between two nodes could be represented in either direction.
- D) This is never a good idea.

Points: 0 out of 1
Correct answer: A)

Question 60 of 80

In a Neo4j database, which of the below best describes what Nodes are used for?

- **A) Used to represent entities and complex value types in the graph.**
- B) As a table structure that identifies like entities and groups them together.
- C) As a reference holder for keys, which also store values, and foreign key links to other nodes.
- D) As endpoints on either side of a relationship, used to bring together relationships in the graph.

Points: 0 out of 1
Correct answer: A)

Question 61 of 80

What information can we **not** find out by executing the `dbms.listQueries()` procedure?

- A) Bytes allocated for the executing query
- B) Planner used by the query
- C) Username of the user who is executing the query.
- **D) IP address of the machine that is executing the query**

Points: 0 out of 1
Correct answer: D)

Question 62 of 80

Neo4j requires which of the following?

- **A) Java runtime environment.**
- B) Microsoft .NET environment.
- C) A J2EE container server.
- D) A SAN storage system.
- E) None of the above.

Points: 0 out of 1
Correct answer: A)

Neo4j requires a Java runtime environment for deployment.

Question 63 of 80

How many nodes can a single relationship connect?

- A) Exactly two.
- B) Exactly one.
- **C) One or two.**
- D) More than two.
- E) None of the above.

Points: 0 out of 1

Correct answer: C)

A node can have a relationship referencing itself or another node.

Question 64 of 80

Of the following, which would be recommended approaches for tuning and potentially improving performance of Neo4j?

- **A) Optimizing the cache settings so more of the graph fits into memory.**
- B) Increasing the size of the performance global area to increase the number of hits in the cache.
- **C) Distributing reads across a cluster of Neo4j instances for higher concurrent access.**
- D) Increasing the parallel_servers configuration setting so queries can be broken down and worked on by multiple threads.

Points: 0 out of 2

Correct answer: A) C)

Question 65 of 80

Neo4j is available with both open-source and commercial licenses.

- **A) True**
- B) False

Points: 0 out of 1

Correct answer: A)

Neo4j can be used with both the open-source license or a commercial license which includes technical support.

Question 66 of 80

The Raft log only contains transaction data

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)
The Raft log stores other data including cluster membership.

Question 67 of 80

Neo4j version 2.2 introduced basic authentication with the server, which when enabled requires a username and password to be supplied to query the server.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 68 of 80

Which of the following schema objects does Neo4j include?

- A) Tables.
- B) Columns.
- C) Materialized views.
- D) Rows.
- **E) None of the above.**

Points: 0 out of 1
Correct answer: E)
Neo4j is a graph database and does not enforce relational schema objects.

Question 69 of 80

Online backups allow for backing up a running instance of Neo4j, which is only available with the Enterprise Edition.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 70 of 80

Foreign keys are necessary in a graph database. These allow you to determine which nodes are related.

- A) True
- **B) False**

Points: 0 out of 1
Correct answer: B)

Nodes are connected by relationships. Foreign keys are not needed.

Question 71 of 80

All nodes with the same label must have the same property keys.

- A) True
- B) **False**

Points: 0 out of 1

Correct answer: B)

Neo4j does not enforce by default that all nodes with the same label must have the same property keys. This is only guaranteed with Property Existence Constraints.

Question 72 of 80

Which log files should I inspect if Neo4j doesn't start up - on a production installation (not Neo4j Desktop).

- A) logs/main.log
- B) /var/log/neo4j.log
- C) **logs/neo4j.log**
- D) **logs/debug.log**
- E) /dev/null

Points: 0 out of 2

Correct answer: C) D)

Question 73 of 80

Neo4j can be embedded directly into your application by importing the libraries and writing directly with the Java APIs.

- A) **True**
- B) False

Points: 0 out of 1

Correct answer: A)

Question 74 of 80

When using a language driver, which forms of transaction can be automatically replayed on failure?

- A) Auto-commit transactions
- B) **Transaction functions**
- C) Explicit transactions

Points: 0 out of 1

Correct answer: B)

Question 75 of 80

Which statement best defines uniqueness constraints in Neo4j?

- A) Neo4j does not support uniqueness constraints.
- B) A rule in the database that ensures a property value is unique among all nodes.
- **C) A rule in the database that ensures a property value is unique for all nodes with a specific label.**
- D) A rule in the database that ensures a node or relationship is unique.

Points: 0 out of 1

Correct answer: C)

Uniqueness constraints are rules in the database that ensure a property value is unique for all nodes with a specific label.

Question 76 of 80

Which of the following Cypher statements will return actors and the directors who directed their movies?

- **A)**

```
MATCH (actor)-[:ACTED_IN]->(movie)<-[:DIRECTED]-(director)
RETURN actor, director
```

- B)

```
MATCH (actor)-[:ACTED_IN]->(movie)
JOIN (movie)<-[:DIRECTED]-(director)
RETURN actor, director
```

- C)

```
MATCH (actor)-[:ACTED_IN]-(movie)
CONNECT (movie)-[d:DIRECTED]-(director)
RETURN actor, director
```

- D)

```
MATCH (actor)-[a:ACTED_IN]->(movie)<-[b:DIRECTED]-(director)
RETURN a, b
```

Points: 0 out of 1

Correct answer: A)

Question 77 of 80

What type of database is Neo4j?

- A) Key-value store.
- B) Document database.
- **C) Graph database.**

- D) Relational database.
- E) Semantic database.

Points: 0 out of 1
Correct answer: C)
Neo4j is a graph database.

Question 78 of 80

What is the result of the following statement: RETURN null <> false

- A) true
- B) false
- **C) null**
- D) none of the above

Points: 0 out of 1
Correct answer: C)

Question 79 of 80

Cypher has a collection of statistics functions that allow you to identify data points such as the maximum and minimum values, standard deviation, and percentiles.

- **A) True**
- B) False

Points: 0 out of 1
Correct answer: A)

Question 80 of 80

When using a language driver, a statement results comprises a stream of...

- A) nodes
- B) rows
- C) entries
- **D) records**

Points: 0 out of 1
Correct answer: D)

Retrieved from 'https://wiki.glitchdata.com/index.php?title=Neo4J:_Sample_Questions_5&oldid=23143'

This page was last modified on 10 October 2018, at 03:45.

Content is available under [Creative Commons Attribution](#) unless otherwise noted.

